

Exercise 1: EM Algorithm for Coin Toss Mixture Model

Problem Setup

We have two coins:

- Coin c_0 : Heads probability q_0
- Coin c_1 : Heads probability q_1

A player picks coin c_1 with probability θ and coin c_0 with probability $1 - \theta$. For each trial i , we observe $x_i \in \{H, T\}$ and have an unobserved latent variable $z_i \in \{0, 1\}$ indicating which coin was chosen. Our goal is to estimate θ using the EM algorithm.

a) Complete-data likelihood $p(x, z|\theta)$

For a single observation (x, z) :

$$p(x, z|\theta) = p(z|\theta) \cdot p(x|z) \quad (1)$$

where:

$$p(z = 1|\theta) = \theta, \quad p(z = 0|\theta) = 1 - \theta$$

If $z = 0$: $p(x = H|z = 0) = q_0$, $p(x = T|z = 0) = 1 - q_0$. If $z = 1$: $p(x = H|z = 1) = q_1$, $p(x = T|z = 1) = 1 - q_1$. Thus:

$$p(x, z|\theta) = \begin{cases} (1 - \theta)q_0^{\mathbb{1}(x=H)}(1 - q_0)^{\mathbb{1}(x=T)}, & z = 0 \\ \theta q_1^{\mathbb{1}(x=H)}(1 - q_1)^{\mathbb{1}(x=T)}, & z = 1 \end{cases} \quad (2)$$

b) Observed-data likelihood $p(x|\theta)$

Marginalizing over z :

$$p(x|\theta) = (1 - \theta)p(x|z = 0) + \theta p(x|z = 1) \quad (3)$$

So:

$$p(x|\theta) = (1 - \theta)[q_0^{\mathbb{1}(x=H)}(1 - q_0)^{\mathbb{1}(x=T)}] + \theta[q_1^{\mathbb{1}(x=H)}(1 - q_1)^{\mathbb{1}(x=T)}] \quad (4)$$

This is a mixture model (Bernoulli mixture), not a Gaussian mixture.

c) Missing-data distribution $p(z|x, \theta)$

By Bayes' theorem:

$$p(z=1|x, \theta) = \frac{\theta p(x|z=1)}{p(x|\theta)} \quad (5)$$

So:

$$p(z=1|x, \theta) = \frac{\theta q_1^{\mathbb{1}(x=H)}(1-q_1)^{\mathbb{1}(x=T)}}{(1-\theta)q_0^{\mathbb{1}(x=H)}(1-q_0)^{\mathbb{1}(x=T)} + \theta q_1^{\mathbb{1}(x=H)}(1-q_1)^{\mathbb{1}(x=T)}} \quad (6)$$

Similarly $p(z=0|x, \theta) = 1 - \gamma(x)$.

d) Q-function for EM

For N observations:

$$Q(\theta|\theta^{(t)}) = \mathbb{E}_{Z|X,\theta^{(t)}}[\log p(X, Z|\theta)] \quad (7)$$

Complete-data log-likelihood:

$$\log p(X, Z|\theta) = \sum_{i=1}^N [z_i \log \theta + (1-z_i) \log(1-\theta)] + \text{const} \quad (8)$$

So:

$$Q(\theta|\theta^{(t)}) = \sum_{i=1}^N [\mathbb{E}[z_i] \log \theta + (1 - \mathbb{E}[z_i]) \log(1-\theta)] \quad (9)$$

Where $\mathbb{E}[z_i] = p(z_i=1|x_i, \theta^{(t)}) = \gamma_i$.

E-step

Compute:

$$\gamma_i = p(z_i=1|x_i, \theta^{(t)}) \quad (10)$$

e) M-step

Maximize:

$$Q(\theta) = \sum_{i=1}^N [\gamma_i \log \theta + (1 - \gamma_i) \log(1-\theta)] \quad (11)$$

Derivative:

$$\frac{\partial Q}{\partial \theta} = \sum_{i=1}^N \frac{\gamma_i}{\theta} - \sum_{i=1}^N \frac{1 - \gamma_i}{1 - \theta} = 0 \quad (12)$$

Solve:

$$\theta^{(t+1)} = \frac{\sum_{i=1}^N \gamma_i}{N} \quad (13)$$

Final EM Algorithm

1. Initialize $\theta^{(0)}$.
2. **E-step:** Compute γ_i for each observation.
3. **M-step:** Update $\theta^{(t+1)} = \frac{1}{N} \sum_{i=1}^N \gamma_i$.
4. Repeat until convergence.